

Improving Keyphrase Extraction with Data Augmentation and Information Filtering^{*}

Amir Pouran Ben Veyseh^{1,*}, Nicole Meister², Franck Dernoncourt³ and Thien Huu Nguyen¹

¹*Dept. of Computer and Information Science, University of Oregon, Eugene, OR, USA*

²*Department of Electrical and Computer Engineering, Princeton University*

³*Adobe Research, Seattle, WA, USA*

Abstract

Keyphrase extraction is one of the essential tasks for document understanding in NLP. While the majority of the prior works are dedicated to the formal setting, e.g., books, news, or web-blogs, informal texts such as video transcripts are less explored. To address this limitation, in this work, we present a novel corpus and method for keyphrase extraction from the transcripts of the videos streamed on the Behance platform. More specifically, in this work, a novel data augmentation is proposed to enrich the model with background knowledge about the keyphrase extraction task from other domains. Extensive experiments on the proposed dataset show the effectiveness of the introduced method.

Keywords

Key Phrase Extraction, Video Transcripts

1. Introduction

Keyphrases are one or different continuous words that may speak to the most thoughts in a report. Keyphrases are commonly categorized as Present or Absent. A present keyphrase unequivocally shows up within the document, while an Absent keyphrase does not exist within the record. Keyphrases can serve as brief rundown for a archive, thus profiting different NLP applications Information Recover [1] and Content Summarization [2]. Due to their value, within the more than two decades, KP has been considered in numerous inquire about works [3, 4, 5, 6, 7, 8, 9].

Though customarily highlight designing has been utilized for KP [3, 10], as of late profound learning is demonstrated to be more productive for this errand [9, 11]. Be that as it may, one restriction within the current works is that they are basically restricted to the formal content such as logical papers [12] and web-logs [13]. As such, the challenges in other spaces are still uncertain. Among others, video transcript is one of the less-explored spaces that may altogether advantage from KP. For occurrence, it may be utilized for video summarization and

VTU'22: The AAAI-2022 Workshop On Video Transcript Understanding, Feb 28, 2022, Virtual

^{*}You can use this document as the template for preparing your publication. We recommend using the latest version of the ceurart style.

✉ apouranb@cs.uoregon.edu (A. P. B. Veyseh); nmeister@alumni.princeton.edu (N. Meister); franck.dernoncourt@adobe.com (F. Dernoncourt); thien@cs.uoregon.edu (T. H. Nguyen)



© 2022 Copyright for this paper by its authors. Use permitted under Creative Commons License Attribution 4.0 International (CC BY 4.0).

CEUR Workshop Proceedings (CEUR-WS.org)

So. But if you guys want to paint the. Image that I'm working on I'm going to share the link. It's actually it should be already available. If you look under the info tab. I believe is above the chat you should be able to get the image. But just in case I want to make extra convenient under the put the link in chat right now, so that should take you through this **picture** if you want to **grab** it and give it ago. I'm going to be **painting** between 3:00 and 4:00 hours on this today, but probably close to 3 hours. My goal is 3 but if it goes over a little bit, then that's OK.

We'll see. We'll see what it what we need to do to make it look good, and we're just going to do that, so that's my plan, but is close to furious I can an. I'm going to put this on my other monitor in **larger size** just like you saw on the **screen** as usual. Are you sick your voice seems little weaker really? I'm not sick? I hope I don't feel sick? I am a little tired getting back from traveling over the weekend and then having like a lot of things going on instantly so I haven't been getting as much sleep as I would like but, yeah, so maybe that's doing it, but no, I I hope I'm not getting sick. I definitely don't need that. Please No.

Wish I could join have an interview thing have to go in a bit might do. It later really like this reference OK awesome. Good luck with your interview and everything. Hi Tesla welcome. Let's put this reference up here using pure ref so it's going to be above are canvas. You can take a look at it no matter what and my **canvas size** if you're curious is 3000 by 2250 pixels. So that's my canvas size here. OK, so I'm not going to have a timer on the. Screen but will just you know **stream** starts around 12:30 so. Around my time, 3:30 or 4 that's kind of what we're going to go for for the **amount of time** on this painting. So. Let's begin I I really love this image when I first saw the references or the images that Andy shared

Figure 1: Part of the transcript of a live-streamed video. The paragraphs are separated by dashed lines and the keyphrases are shown in red boldface.

recovery or advantage individuals who are hard of hearing and difficult of hearing (DHH) [14]. On the other hand, KP for transcripts that are consequently gotten are more challenging than the formal composed records as these transcripts include loud content, incomplete/repeated sentences and expressions, casual lexicon, and non cohesive data stream. In spite of the fact that there have been many related endeavors to assess highlight designing strategies on assembly transcripts [10, 15], the accessible assets, with a modest bunch of transcripts and keyphrases, are not valuable to train/evaluate the later progressed profound models.

To address these limitations, in this work, we study the task of keyphrase extraction in the domain of video transcripts. In particular, we explore the transcript of the videos streamed on the Behance¹. An example of these transcripts are presented in Figure 1. To this end, first a collection of the video transcripts are annotated with keyphrases. Next, we explore deep learning models to extract the keyphrases from the the transcripts of the videos. One of the challenges for the keyphrase extraction in video transcripts is that the amount of training samples in this domain is less than the well-studied domains such as news. As such, in a novel approach, in this work we present a data augmentation technique, in which the data from other domains is employed to improve the performance of the keyphrase extraction model. Our extensive analysis on the proposed dataset reveal the effectiveness of the proposed approach.

2. Related work

This task could be modeled as extractive keyphrase extraction (EKE) [16] whose goal is to identify the salient phrases in a document. These systems are designed to identify the word(s)

¹www.behance.net

in a document that forms a phrase and refer to key points in the document. However, there are some differences that render the existing methods for extractive keyphrase extraction inapplicable for our task. First, EKE is conducted on the formal text where the sentences are grammatically correct and no chitchat or repeated information exists in the text. In contrast, in the proposed task we deal with automatically transcribed videos that might be noisy and also it might have redundant information in the forms of chitchat or repeated sentences/words. As such, since EKE systems are not designed to avoid uninformative portions of the input, they will fail on this task. Second, one of the requirements in the proposed task is to ensure the uniqueness of the keyphrases of each paragraph. Existing EKE systems are not designed to observe this requirement, thereby they might extract the same phrase for multiple paragraphs. Third, the existing EKE systems are trained in the scientific or weblogs domain. They are substantially different from the domain of live-stream videos such as Behance videos that generally contain technical content. Moreover, none of the prior works can benefit from the available resources in the general domain (i.e., weblogs) for a domain-specific EKE system. As such, in this work, we propose a novel system that can work on noisy live-stream video transcripts meanwhile benefit from available resources in the general domain. Our proposed system is the first model that is encouraged to ensure the uniqueness of the keyphrases.

3. Model

3.1. Overview

Our proposed model has the following novelties:

- A novel method for extracting keyphrases from the paragraphs of a live-stream video
- A novel technique to encourage the uniqueness of the keyphrases across consecutive paragraphs of the transcript
- A novel method for identifying the domain-specific phrases using a domain discriminator model
- A novel method based on multi-tasking to bridge the phraseness and keyness information from general domain to a specific domain
- A novel method based on reinforcement learning to inform the model about the chitchats in the transcript

3.2. Details

Formally, the input to the system is a paragraph of a transcript, i.e., $D = [w_1, w_2, \dots, w_n]$, consisting of n words. The goal is to select phrases P in D that most clearly represent the main content of the paragraph D . A phrase P might be a single word or a span in D . All noun, verb, or adjective phrases are eligible to be selected as one of the phrases P . In this work, the phrases P are encoded in the label sequence $L^P = [l_1^P, l_2^P, \dots, l_n^P]$, where $l_i^P \in \{O, B, I\}$ and B indicates the word w_i is the beginning of a phrase and I indicates the word w_i is the continuation of a phrase. To create a system for this task, we propose multiple components. Specifically, the proposed system consists of three major components:

- **Keyphrase Extractor:** The input paragraph D is first encoded to high-dimensional vectors $H = [h_1, h_2, \dots, h_n]$ using a pre-trained transformer-based language model. Moreover, to preserve the information about the phrases P' extracted from the previous paragraph D' , these phrases are also concatenated to the input paragraph D . Using the extracted feature vectors for the input words, we predict the label for each word.
- **Data Augmentation:** In addition to the training samples for keyphrase extraction from transcripts, we propose to employ the general domain training samples for keyphrases extraction. However, the domain shift between these resources could impede the training. As such, we devise a component dedicated to bridging the gap between general domain and transcript domain keyphrases. Specifically, this component aims to transfer the keyness and phraseness features learned from the general domain to the domain of transcripts. It is achieved by predicting the bridge phrases. In particular, in addition to the main keyphrase extraction, we suggest training the model to recognize phrases that are representative of their domain (e.g., transcripts or general domain). This task could help the model to get informed of the portions of the input that are specific to the domain of interest, thereby filter out irrelevant knowledge/features related to the other domain. We propose a novel method to obtain a pseudo-label for this task. Specifically, a separate transformer-based language model is trained to discriminate the domains of the paragraphs and the attention scores for each phrase from the discriminator model are employed to construct domain-specific phrases.
- **Information Filter:** In this task, we require the keyphrases to be selected from the informative portion of the input paragraph. That is, chitchats and other non-informative portions should be avoided. In order to encourage the model to observe this requirement, we first create pseudo-binary labels for each sentence of the paragraph where 0 indicates the sentence is chitchat and 1 indicates the sentence is informative. Next, we reward the keyphrase extraction model if the selected keyphrases are from informative sentences.

The rest of this section elaborates more on the details of each of these components.

3.2.1. Keyphrase Extraction

In this work we use BERT_{base} [17] as the input encoder. It is worth noting that in addition to the input paragraph, we suggest concatenating the keyphrases of the previous paragraph to the input text so that the encoder is aware of the words that it will later be encouraged to avoid selecting them as keyphrases. To this end, to encode the input paragraph D and the keyphrases of the previous paragraph, we construct the sequence $S = [w_1, w_2, \dots, w_n, SEP, kp_1, kp_2, \dots, kp_m]$, where $kp_i \in P'$ is the i -th keyphrase extracted for the previous paragraph. To find kp 's, during training we use the gold keyphrases and during inference, we use the keyphrases predicted by the model. The sequence S is fed into the BERT_{base} model and the representations of the word-pieces of the input paragraph is taken from the final layer of the BERT transformer as the input word representations $H = [h_1, h_2, \dots, h_n]$. Note that for the words consisting of multiple word pieces we use the average of their word-piece representation to obtain their final vector.

Next, using the vectors H , we predict the likelihood of every word to be selected in a keyphrase. Formally, the vector h_i is consumed by a two-layer feed-forward network to estimate the likelihood $P(\cdot|D, w_i, \theta)$:

$$P(\cdot|D, w_i, \theta) = \sigma(W_2(W_1 * h_i + b_1) + b_2, \theta) \quad (1)$$

where σ is the softmax function, W_2 and W_1 is the weight matrices, b_1 and b_2 are the biases, and $P(\cdot|D, w_i, \theta)$ is the label distribution for the word w_i predicted by the model with parameters θ . To train the model for keyphrase extraction, we use the following cross-entropy loss:

$$\mathcal{L}_{kp} = - \sum_{i=1}^n \log(P(l_i|D, w_i, \theta)) \quad (2)$$

Also, in order to encourage the model to avoid selecting repeated keyphrases for the consecutive paragraphs, we compute the following reward:

$$R_{rep}(KP) = -\frac{1}{n} \sum_{i=1}^n \text{REP}(w_i) \quad (3)$$

$$\text{REP}(w_i) = \begin{cases} 1, & \text{if } w_i \in P' \quad \& \\ & \text{argmax}(P(\cdot|D, w_i, \theta)) \in \{B, I\} \\ 0, & \text{otherwise} \end{cases}$$

where KP is the list of predicted keyphrases and $\text{REP}(w_i)$ is a function that returns 1 if the word w_i is predicted to be in a keyphrase and also it appears in the keyphrases of the previous paragraph, i.e., P' .

3.2.2. Data Augmentation

In this work, we extract keyphrases from live-stream video transcripts. However, there are other manually annotated resources in other domains which might be helpful to improve the performance of the keyphrase extraction model. One of these resources is OpenKP [13] that provides human annotation for keyphrase extraction in the domain of web pages. This dataset provides annotation for about seventy thousand web pages. One simple method to employ this resource in the training of the model is to combine this dataset with the annotated samples of the live-stream videos. While this simple technique could help the model to learn more patterns of the keyphrases, one limitation is that the domain shift between web pages and live-stream videos might hinder the model training. Specifically, a model trained on the web page domain might pay more attention to chitchats of the live-stream videos and ignores the informative portion that contains more technical phrases. As such, it is necessary to equip the model with a mechanism to overcome the domain shift. In this work, we suggest that identifying the keyphrases that are representative of their domain is critical information that a model trained on multiple domains should be aware of in order to avoid selecting keyphrases that are suitable for the other auxiliary domain. In other words, we aim to train the model in a multitask setting, in which in addition to the main keyphrase extraction, the model is trained to recognize domain-specific phrases. Therefore, the first step is to create labels for the phrases that are domain-specific in the paragraphs of the live-stream video and the web pages.

Domain-specific Phrase Annotation: The purpose of domain-specific phrase detection is to recognize the phrases $P' \in D$ that are representative of the domain of D . Since there

is no labeled data for this task, we resort to an unsupervised method. In the first step, we suggest automatically construct a labeled dataset for domain-specific phrase detection using a pre-trained domain discriminator. In particular, we first combine all training paragraphs D from the live-stream video domain and D' from the web page domain to construct the dataset \mathcal{D} . Next, we employ a BERT_{base} model² to encode the paragraph $\bar{D} \in \mathcal{D}$. The representations $\bar{H} = [\bar{h}_1, \bar{h}_2, \dots, \bar{h}_n]$ obtained from the last layer of the BERT_{base} model³ are later max-pooled (mp()) and are consumed by a feed-forward layer to predict the domain of \bar{D} :

$$\begin{aligned} V &= mp(\bar{h}_1, \bar{h}_2, \dots, \bar{h}_n) \\ P(\cdot|\bar{D}, \theta') &= \sigma(W_2 * (W_1 * V + b_1) + b_2) \end{aligned} \quad (4)$$

where σ is the sigmoid activation function, W_1 and W_2 are the weight matrices, b_1 and b_2 are the biases and $P(\cdot|\bar{D}, \theta)$ is the probability distribution over the two domain (i.e., live-stream videos and web pages) predicted by the discriminator with parameter θ' . This pre-trained model has an accuracy of 93% on the test set of \mathcal{D} .

Next, we apply the pre-trained domain discriminator on all documents $\bar{D} \in \mathcal{D}$ to obtain the attention scores $A = [a_1, a_2, \dots, a_n]$ for all words $w_i \in \bar{D}$. Note that these attention scores are obtained from the final layer of BERT_{base} encoder of the discriminator. Using the attention scores A , we filter the words of the document $\bar{D} \in \mathcal{D}$. Specifically, in the first step, the attention scores A are sorted descendingly, i.e., $A' = [a'_i | a'_i \in A \quad \& \quad A'_i \geq A'_j \Rightarrow i \leq j]$. We define the function $g(x)$ whose input is the index i of a'_i and its output is the index j of a_j corresponding to entry a'_i , i.e. $a_j = a'_i$. Next, the array W' is constructed whose entry $w'_i \in W'$ is the word $w_j \in \bar{D}$ where $j = g(i)$, i.e., the word corresponding to the i -th entry of sorted attention score A' . Finally, we prune all words of the paragraph \bar{D} that appear after the index k in W' :

$$\hat{\bar{D}} = \{w_i | w_i \in \bar{D} \quad \& \quad \text{Index_of}(w_i, W') < k\} \quad (5)$$

where $\text{Index_of}(x, Y)$ is a function that returns the index of x in the array Y . In order to find the optimal value for k , we use the following criteria:

$$|P(\cdot|\hat{\bar{D}}, \theta') - P(\cdot|\bar{D}, \theta')| \leq \eta \quad (6)$$

where $P(\cdot|\hat{\bar{D}}, \theta')$ is the domain distribution predicted by the pre-trained discriminator for the filtered document $\hat{\bar{D}}$, and η is a threshold to be selected based on the performance on the development set. The main motivation for this filtering criteria is that only those important words that are necessary to make the same prediction as the original input document should be preserved in the filtered document.

Every word that is remained in the filter document $\hat{\bar{D}}$ is used to create the silver data for domain-specific phrase detection. Concretely, the label vector $L = [l_1, l_2, \dots, l_n]$ is constructed as the silver labels for domain-specific phrase detection:

²Note that it is separate from the main model encoder

³We compute the average vector representation for the words with multiple word-pieces

$$l_i = \begin{cases} 1, & \text{if } w_i \in \hat{D} \\ 0, & \text{otherwise} \end{cases} \quad (7)$$

Finally, using the silver data, the main keyphrase extraction model is trained to recognize the domain-specific phrases in a multi-task setting. In particular, the vector representation $H = [h_1, h_2, \dots, h_n]$, obtained from the keyphrase extraction encoder, are fed into a feed-forward layer with sigmoid activation function at the end to predict the domain-specific phrases:

$$Q(\cdot|D, w_i, \tilde{\theta}) = \sigma(W_2(W_1 * h_i + b_1) + b_2) \quad (8)$$

where σ is the sigmoid activation function, W_1 and W_2 are weight matrices, b_1 and b_2 are biases and $Q(\cdot|D, w_i, \tilde{\theta})$ is the distribution of the labels (i.e., domain-specific or not), predicted by the model with parameter $\tilde{\theta}$ for i -th word. To train the model using the silver labels L , we use the following negative log-likelihood:

$$\mathcal{L}_{bridge} = - \sum_{i=1}^n \log(Q(l_i|D, w_i, \tilde{\theta})) \quad (9)$$

3.2.3. Information Filter

In a live-stream video, the streamer might diverge from the main content of the video. These portions of the transcripts, which we call chitchats, should be avoided for keyphrase extraction as they are not informative. To this end, we add another component to our model which encourages the model to avoid selecting keyphrases from chitchats. In this work, we propose to do it at the sentence level. That is, keyphrases should not be selected from sentences that are likely to be chitchat. Now, the main question is how can we effectively determine which sentences are chitchat. To answer this question, we resort to an unsupervised method based on the semantic representation of the sentences of the paragraph D . Specifically, we hypothesize that the majority of the sentences in the paragraph are on the topic sentences. As such, in order to detect the sentences that are off the topic, we propose to compute the semantic similarity of the sentences with the entire paragraph. Those sentences whose representations are far enough from the paragraph representation are selected as the chitchat sentences. Formally, given the document D encoded by the BERT_{base} model, i.e., the vectors H , we take the $[CLS]$ vector representation obtained from the final layer of the BERT transformer as the document level representation, i.e., h_p . Next, the representation of the sentence $S_i \in D$ is computed via the max-pooling over its word vector representations:

$$h_{S,i} = \text{mp}(\{h_k | w_k \in S_i\}) \quad (10)$$

Afterward, using the paragraph representation h_p and the sentence representations $h_{S,i}$, we compute a score for every sentence S_i :

$$\alpha_i = \sigma(h_p) \odot \sigma(h_{S,i}) \quad (11)$$

where σ is the softmax operation and \odot is the Hadamard product. Finally, we choose the chitchat sentences based on their computed scores α_i :

$$\text{Is_Chitchat}(S_i) = \begin{cases} 1, & \text{if } \alpha_i \leq \beta \\ 0, & \text{otherwise} \end{cases} \quad (12)$$

where β is a trade-off parameter to be selected based on the performance on the development set. Finally, using the selected chitchat, we define the following reward function:

$$R_{chitchat}(KP) = - \sum_{i=1} \text{Is_Chitchat}(\text{Sent}(kp_i)) \quad (13)$$

where $kp_i \in KP$ is the i -th keyphrase selected by the model and $\text{Sent}(x)$ is a function that returns the sentence containing keyphrase x .

3.2.4. Training

To train the model on the entire task we combined the keyphrase extraction loss \mathcal{L}_{kp} , the bridge loss \mathcal{L}_{bridge} and the rewards $R_{rep}(D)$ and $R_{chitchat}(D)$. Since the reward computation is a discrete operation, we resort to Reinforce algorithm to compute the gradients of the rewards. First, the overall reward is computed by $R(KP) = R_{rep}(KP) + \alpha R_{chitchat}(KP)$, where α is a trade-off hyper-parameter. Next, we seek to minimize the negative expected reward $R(KP)$ over the possible choices of KP : $\mathcal{L}_R = -\mathbb{E}_{\hat{KP} \sim P(\hat{KP}|D)}[R(\hat{KP})]$. The policy gradient is then estimated by: $\nabla \mathcal{L}_R = -\mathbb{E}_{\hat{KP} \sim P(\hat{KP}|D)}[(R(\hat{KP}) - b)\nabla \log P(\hat{KP}|D)]$. Using one roll-out sample, we further estimate $\nabla \mathcal{L}_R$ via the predicted keyphrases KP : $\nabla \mathcal{L}_R = -(R(KP) - b)\nabla \log P(\hat{KP}|D)$ where b is the baseline to reduce variance. In this work, we obtain the baseline b via: $b = \frac{1}{|B|} \sum_{i=1}^{|B|} R(KP_i)$, where $|B|$ is the mini-batch size and KP_i is the predicted keyphrases for the i -th sample in the mini-batch.

3.3. Dataset

To train and evaluate the model we annotate data from the transcripts of the videos streamed on the Behance platform. The recordings are spilled by specialists and creators to share/discuss their inventive ventures. As such, verbal substance from the speakers (in English) is imperative for video understanding. Whereas the recordings have introductory subjects, their substance is impromptu, thus the streamer might cut sentences, examine numerous themes, and utilize casual expressions. The recordings have a normal length of 48 minutes. To get the verbal substance of the gushed recordings, we utilize the Microsoft ASR tool. In add up to, 361 recordings with a add up to length of more than 500 hours are transcribed. A transcript, on normal, contains 7,219 words.

As discussed in the introduction, the long nature of transcripts spurs us to explain keyphrases at two levels. To begin with, at the section level, we characterize a passage in a transcript to have the same part as passages in formally composed archives. Concretely, a passage is characterized as a chunk of content that passes on a specific point or thought. A transcript comprises of numerous disjoint passages. Since the ASR content does not give section data, we clarify the collected transcripts with sections. Next, for each section of the transcript, the critical keyphrases are chosen. To this conclusion, a keyphrase for a passage ought to have

Model	F1@1	F1@3	F1@5
JointKPE (BERT)	14.44	18.91	24.19
JointKPE (RoBERTa)	16.00	22.07	25.08
JointKPE (SpanBERT)	16.08	24.96	27.63
Ours	28.50	36.43	33.83

Table 1
Performance of the models.

these features: (a) Concisely summarize the most idea within the section; (b) Be related to the most subject of the video; (c) Expressly show up within the passage; (d) Does not appear within the past or following passages; (e) Shape a legitimate English noun/verb state. The passages that are totally off-topic don’t have any keyphrase. Next, at the chapter level, we offer keyphrases for chapters within the transcripts. A chapter comprises of numerous sections to speak to a single point. For instance, in a photo altering video, the talk on how to alter the foundation can shape a chapter. A keyphrase of a chapter should observe the following criteria: (a) Concisely summarize the most subjects within the chapter; (b) May not unequivocally show up within the chapter; (c) Does not cover with the section keyphrases or other chapter level keyphrases; (d) Shape a appropriate English noun/verb state. Note that passages and chapters might have different keyphrases that are sorted based on their significance.

To explain information for each level, we enlist 10 annotators from the upwork . com platform which is a site for enlisting freelancers with distinctive mastery. Since the collected recordings are related to photo altering program, e.g., Photoshop, we require the annotators to have encounter both in information explanation and in utilizing major photo altering apparatuses. We prepare the annotators for KP at each level. To anticipate chapter-level keyphrases to be one-sided toward paragraph-level keyphrases, we part annotator pool for section and chapter level explanation (five for each). The transcripts are disseminated equally to the five annotators at each level for comment. As such, a transcript is explained completely by a passage annotator and a chapter annotator.

Figure 2 shows an example of the annotation tool. In this tool the annotator can select/correct the boundaries of the paragraphs in each video transcripts. Each paragraph are shown in different colors. For each paragraph, the annotator selects a group of the keyphrases from the paragraph. In the example shown in Figure 2, the keyphrases “*combination, painting, foreground*” are provided. For more examples of the annotation tool, see appendices.

3.3.1. Results

We evaluate the performance of the keyphrase extraction model on the dataset of live-stream video transcripts. In our experiments, we compare our model with the previous state-of-the-art keyphrase extraction model JointKPE [16]. We compare the models using the F1 score at the first 1, 3, and 5 keyphrases extracted by the model⁴. Table 1 show the results of this comparison. As shown in this table, the proposed model significantly outperforms the baselines in all metrics. This superior performance demonstrates the effectiveness of the proposed model.

⁴Keyphrases are sorted based on the likelihood of their first word, i.e., $P(B|D, w_i, \theta)$

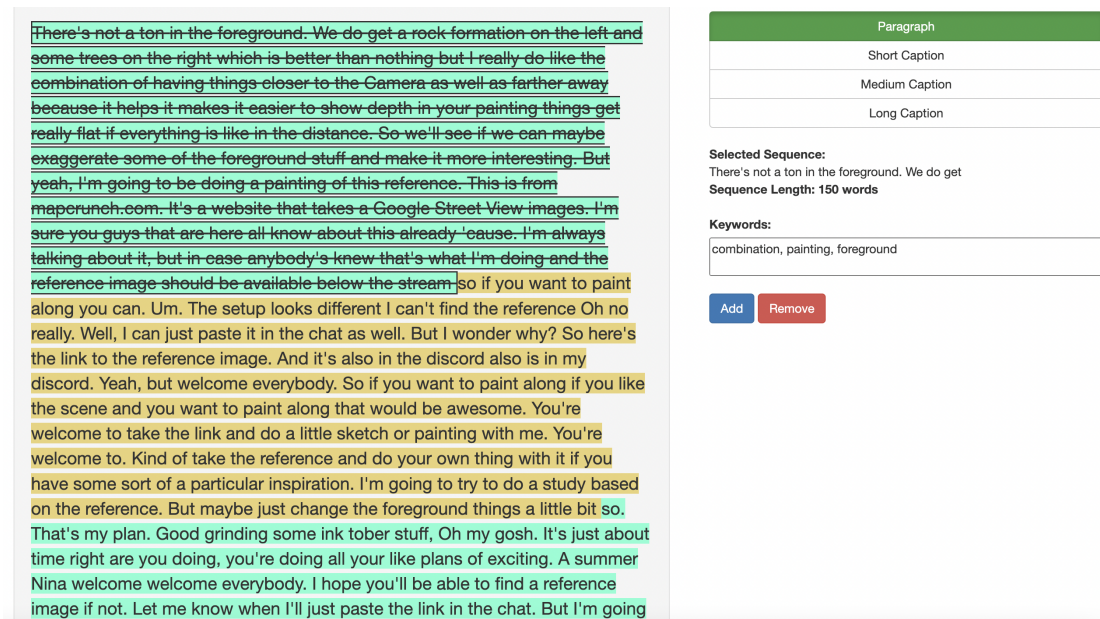


Figure 2: Screenshot of the annotation tool

References

- [1] W. Hersh, Information retrieval, in: Biomedical Informatics, 2021.
- [2] A. Adhikari, A. Ram, R. Tang, W. L. Hamilton, J. Lin, Exploring the limits of simple learners in knowledge distillation for document classification with docbert, in: Proceedings of the 5th Workshop on Representation Learning for NLP, 2020.
- [3] P. D. Turney, Learning algorithms for keyphrase extraction, in: Information retrieval, 2000.
- [4] Y.-f. B. Wu, Q. Li, R. S. Bot, X. Chen, Domain-specific keyphrase extraction, in: Proceedings of the 14th ACM international conference on Information and knowledge management, 2005, pp. 283–284.
- [5] X. Jiang, Y. Hu, H. Li, A ranking approach to keyphrase extraction, in: Proceedings of the 32nd international ACM SIGIR conference on Research and development in information retrieval, 2009.
- [6] K. S. Hasan, V. Ng, Automatic keyphrase extraction: A survey of the state of the art, in: Proceedings of the 52nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), 2014.
- [7] D. Mahata, J. Kuriakose, R. R. Shah, R. Zimmermann, Key2Vec: Automatic ranked keyphrase extraction from scientific articles using phrase embeddings, in: Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 2 (Short Papers), 2018. URL: <https://aclanthology.org/N18-2100>. doi:10.18653/v1/N18-2100.
- [8] W. Chen, H. P. Chan, P. Li, I. King, Exclusive hierarchical decoding for deep keyphrase

generation, in: Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics, 2020. URL: <https://aclanthology.org/2020.acl-main.103>. doi:10.18653/v1/2020.acl-main.103.

- [9] J. Ye, T. Gui, Y. Luo, Y. Xu, Q. Zhang, One2Set: Generating diverse keyphrases as a set, in: Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers), 2021. URL: <https://aclanthology.org/2021.acl-long.354>. doi:10.18653/v1/2021.acl-long.354.
- [10] J. Sheeba, K. Vivekanandan, A fuzzy logic based improved keyword extraction from meeting transcripts, in: International Journal on Computer Science and Engineering, 2014.
- [11] W. Ahmad, X. Bai, S. Lee, K.-W. Chang, Select, extract and generate: Neural keyphrase generation with layer-wise coverage attention, in: Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers), 2021. URL: <https://aclanthology.org/2021.acl-long.111>. doi:10.18653/v1/2021.acl-long.111.
- [12] R. Meng, S. Zhao, S. Han, D. He, P. Brusilovsky, Y. Chi, Deep keyphrase generation, in: Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), 2017. URL: <https://aclanthology.org/P17-1054>. doi:10.18653/v1/P17-1054.
- [13] L. Xiong, C. Hu, C. Xiong, D. Campos, A. Overwijk, Open domain web keyphrase extraction beyond language modeling, in: Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP), Association for Computational Linguistics, Hong Kong, China, 2019, pp. 5175–5184. URL: <https://aclanthology.org/D19-1521>. doi:10.18653/v1/D19-1521.
- [14] S. Kafle, P. Yeung, M. Huenerfauth, Evaluating the benefit of highlighting key words in captions for people who are deaf or hard of hearing, in: The 21st International ACM SIGACCESS Conference on Computers and Accessibility, ASSETS 2019, Pittsburgh, PA, USA, October 28-30, 2019, 2019. URL: <https://doi.org/10.1145/3308561.3353781>. doi:10.1145/3308561.3353781.
- [15] J. Sheeba, K. Vivekanandan, Improved keyword and keyphrase extraction from meeting transcripts, in: International Journal of Computer Applications, 2012.
- [16] S. Sun, C. Xiong, Z. Liu, Z. Liu, J. Bao, Joint keyphrase chunking and salience ranking with bert, arXiv preprint arXiv:2004.13639 (2020).
- [17] J. Devlin, M.-W. Chang, K. Lee, K. Toutanova, BERT: Pre-training of deep bidirectional transformers for language understanding, in: Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers), Association for Computational Linguistics, Minneapolis, Minnesota, 2019, pp. 4171–4186. URL: <https://aclanthology.org/N19-1423>. doi:10.18653/v1/N19-1423.

Paragraph: There's not a ton in the foreground. We do get

Keywords: combination, painting, foreground

Paragraph: so if you want to paint along you can. Um.

Keywords: sketch

Paragraph: so. That's my plan. Good grinding some ink tober stuff,

Keywords: image, link, thumbnail

Paragraph: One is really nice. And we can hopefully get a

Keywords: plans, working, screen, settings

Paragraph: I feel like it's somewhere here mode always on Top

Keywords: brushes, stream, video

Paragraph: The assets are the uploaded assets are visible on my

Keywords: visible, horizon, line

Paragraph: I think at this time it kind of fits it

Keywords: opportunity, the effect

Paragraph: But I think it might be more interesting if there

Keywords: overlapping, fantasy scene, fantasy

Paragraph
Short Caption
Medium Caption
Long Caption

Selected Sequence:
and you want to share what you did, you're welcome
Sequence Length: 130 words

Keywords:

N/A

AddRemove

Figure 3: Annotation tool for providing the keyphrases of the entire transcript

A. Annotation Tool

An example of the annotation tool for providing the summary of the selected keyphrases of the entire transcript is shown in Figure 3. In particular, for each annotated paragraph in the given transcript, the annotators see a list of provided key phrases. Note that in the case that the paragraph is entirely chitchat or off-topic, the annotator provides “N/A” keyphrase.

B. Case Study

To shed more light on the performance of the presented model, in this section, we provide some examples of the extracted keyphrases by the model. Figures 4, 5, and 6 show the examples. In these examples, the predicted keyphrase is shown in red bold-face.

faster so for the few times that it doesn't. And I have to Reselect, something it's OK, but I'm going to make an **adjustment layer**. So it's under my Camera, but there's little half circle icon down there going to press that and go to. Let's try selective color. I think maybe. Maybe I could mess with. Some of the Some some colors on the rock. Might be able to change some things. I when it comes to this, this kind of adjustment. I am just completely winging it. See Select Reselect. Maybe maybe that wasn't the way to go, maybe like curves or levels could have been better. Yeah, there, we go. So lowering the mid tones. That might help with Lower the idea of lowering the contrast a little bit. So I can. Soften the edges of that mask a little bit. And. Yeah. Maybe be like that better. Yeah, I'm at, I'm at maybe it's better like this. I think I think I. I think it's a

Figure 4: Case Study - The keyphrase is shown in red bold-face

things coming up here. i think it would be maybe darker. And then if it comes up above the. That rock. Maybe we get it by the light and have more of like a bright green. Warmer color. Something like that. And maybe we can have some to have different. Type of texture like maybe some individual leaves something more like a Bush rather than. Then. Reeds or I don't know what he called that like. **Pulling a blank**. But you will get a Bush with some larger leaves here. Instead of just like the grass is. A Jackson thank you. Thank you so much welcome do you remember this one I think it was suggested recently? I really like the reference so I'm going to do it now. Body like it, I changed a few things about it not too much, but I wanted to. Do a little bit more with the foreground,

Figure 5: Case Study - The keyphrase is shown in red bold-face

Alright Skip back in the painting. Going to make a new layer above the layer of the. Rocks and | want to start **overlapping** some greenery and and Bush is and things. So it's not just yeah 'cause. It says it feels a little bit fake to me that there's there's this layering and there's nothing coming back on Top so. Let's put a Bush here or or just have a little bit of grass is coming up. I'm just trying to figure out. The texture that | want for that. So. Yeah, I'm going to look at the windows ink thing. After this stream. There's some weird stuff going on still. Maybe we can get rid of that. Where did my brush pen capsules time to try this inktober bays? Enjoy enjoy. It's been too long for me. Too long since I've done ink stuff, Iso really enjoying it used to be my main medium before I got into

Figure 6: Case Study - The keyphrase is shown in red bold-face